

---

# P2: Client/Server File System

Akshay Parashar

Hemal Patel

Kalpita Munot

Reetuparna Mukherjee

---

# Protocol Semantics

# POSIX Implementation



## Local only calls:

- read()
- write()
- utimens()
- flush()

## Local and remote calls:

- getattr()
- readdir()
- mkdir()
- rmdir()
- creat()
- unlink()
- open()
- release()

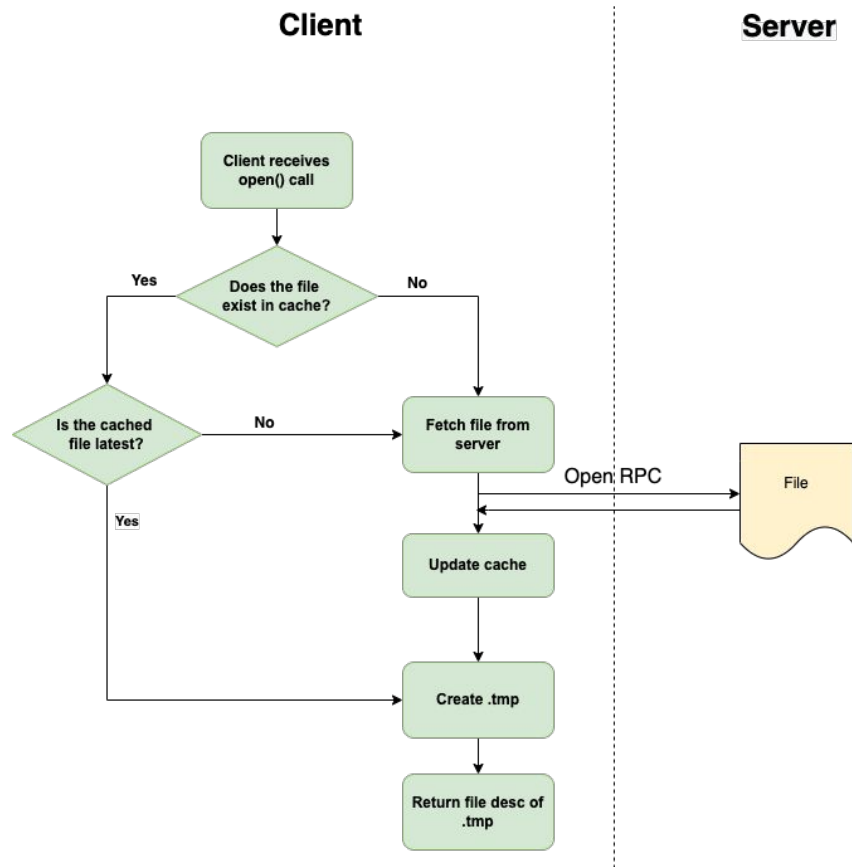
# Terminology and Protocol Assumptions



- **Nomenclature:**
  - **.tmp file:** Dirty files that have not been closed by the application process
  - **.consistent file:** files that are closed by application process and are safe to be flushed on client/server
- **Assumptions:**
  - Unless close() API is called, data is non-recoverable (dirty).
  - After close() is called and if server or client crashes, the data needs to be recovered by crash consistency protocol and updated on both the server and client
  - "Last Writer Wins" on both and Server and Client.

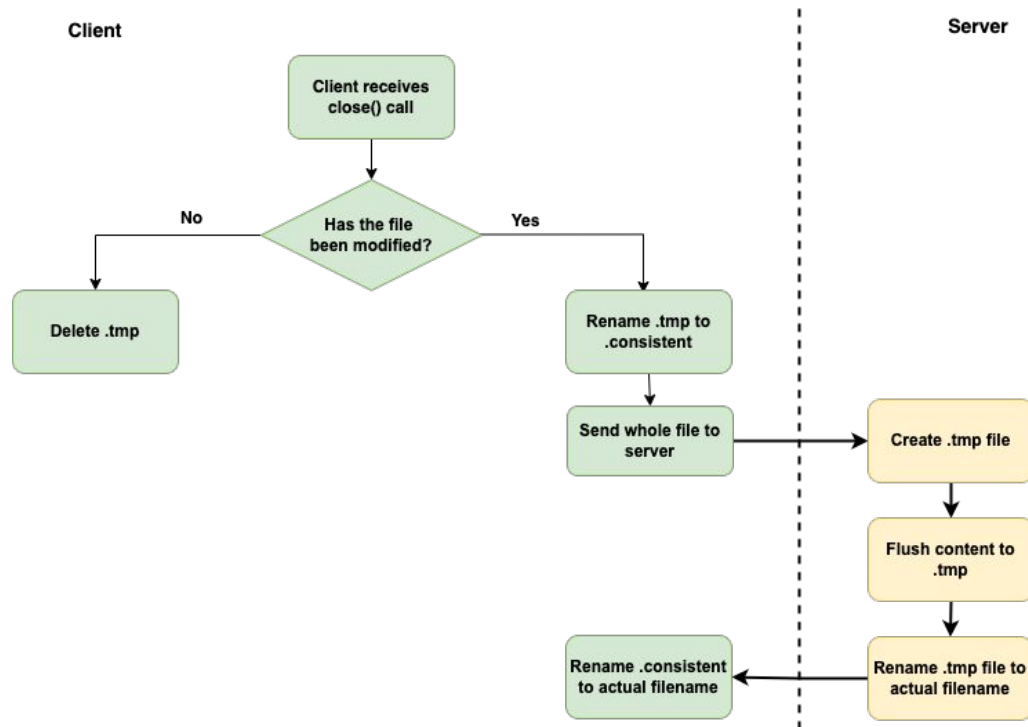
# Open

- Client fetches the file on demand from the server iff
  - file does not exist on cache (or)
  - cached file is stale
- On fetching the file from server, client updates its cache
- Each application process is given its own .tmp file (i.e. copy of the cached file) descriptor to work on



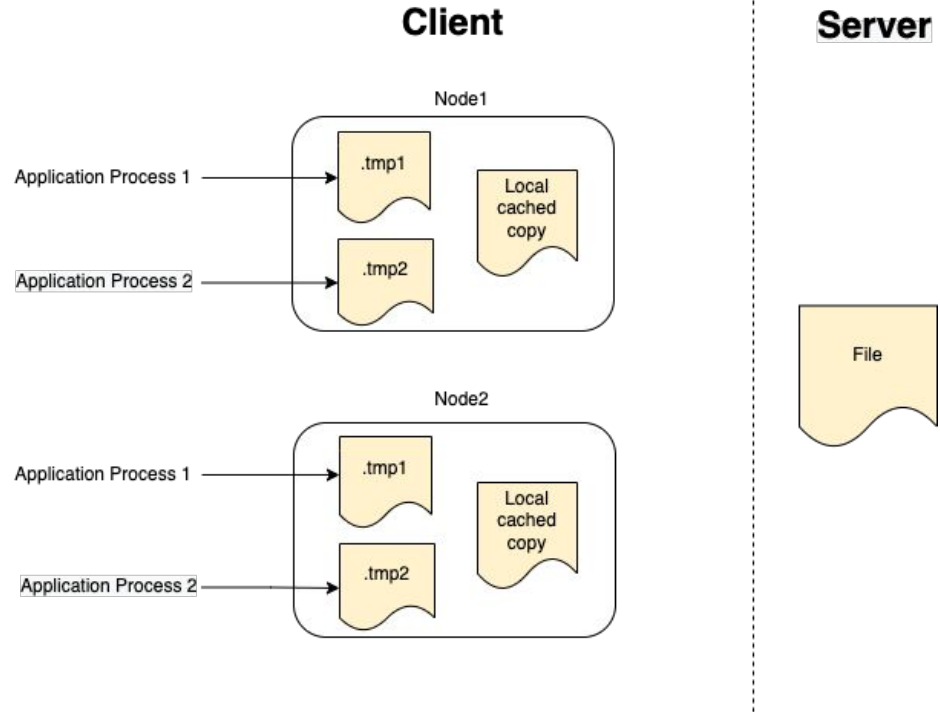
# Close

- Fuse-client checks for changes in .tmp file.
- .tmp file is deleted if there are no changes,
- Else, .tmp file state is changed to .consistent file (renamed)
- File is then sent to the server
- Server creates a .tmp for every close op, and then atomically renames it - LAST WRITE WINS
- After sending the file, .consistent file is used to create the cached copy for that node



# Read/Write

- Read/ Write are local to each client
- Each application process performs read/write on its own copy of cached file (.tmp file created during open)
- No server interaction involved



---

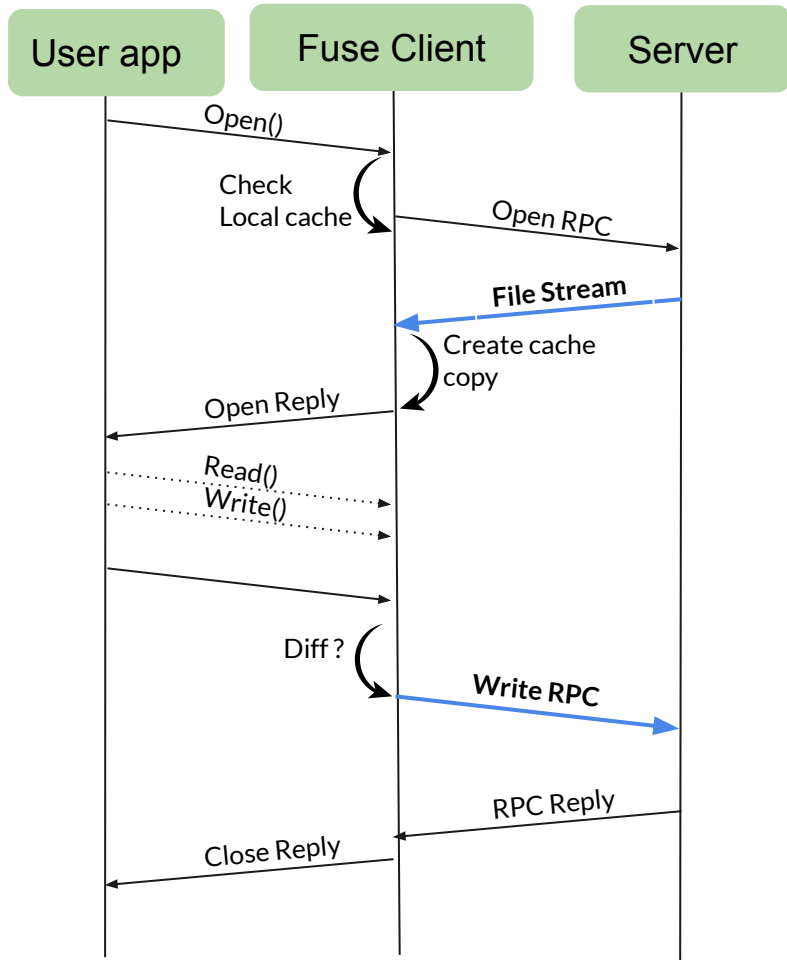
# Durability and Reliability



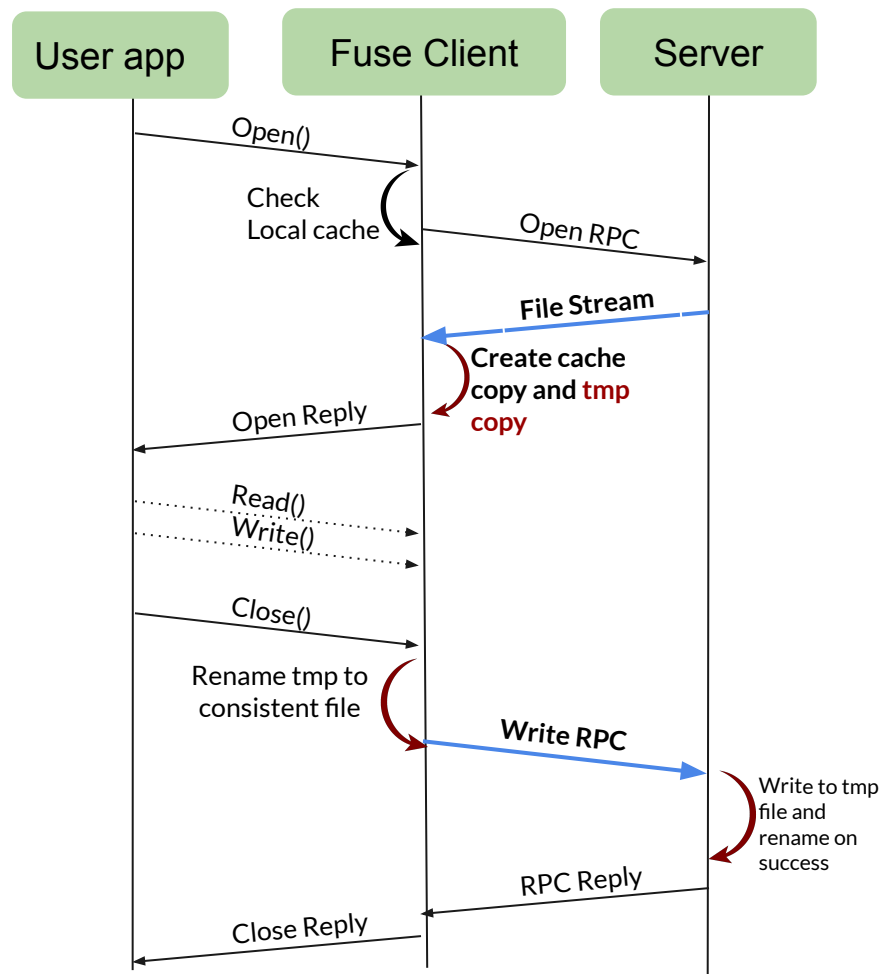
# Durability: Crash Consistency Protocol

- Each application process per client creates a **.tmp** copy for their read/writes
- If the client/server crashes before application process could call close, the tmp file is dirty and is deleted during **garbage collection**
- Renaming .tmp file to **.consistent** file on close saves state of the file from dirty to consistent. It is an **atomic** operation, hence safe
- On restart, client process deletes all the .tmp files and recovers .consistent files to local cache and server.

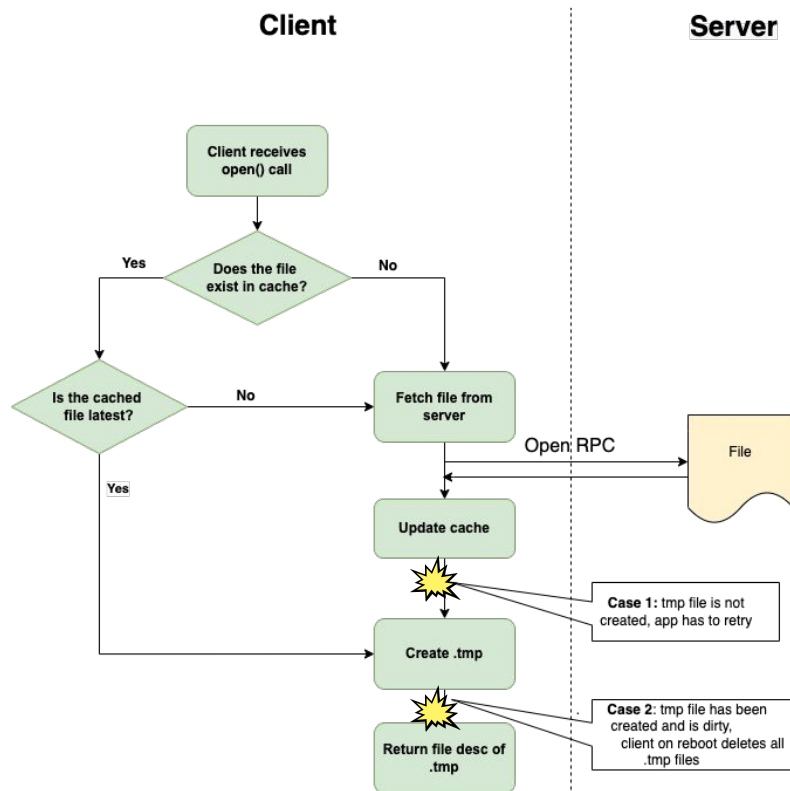
## W/o crash and concurrency



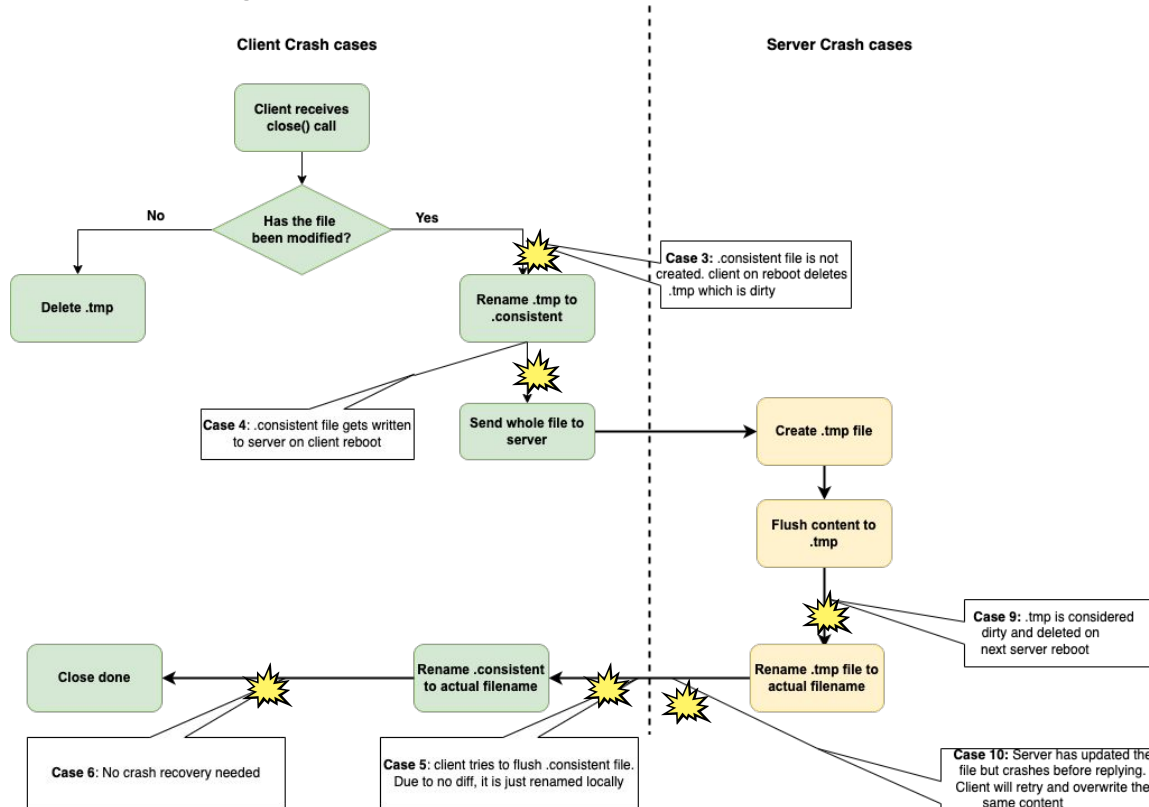
## Our Implementation



# Crash Recovery Protocol, Open



# Crash Recovery Protocol, Close



---

# Performance

# Setup



**Workload:** File size ranges from 2KB to 1024 MB

**Processes count:** 2,4,8,16

## **Hardware Configuration:**

- Cloud Provider: MSFT Azure
- OS: Linux (ubuntu 20.04)
- vCPU: 2
- RAM: 4GB

# Cases

Number	Case description		Calls involved
1	Client creates the file for the first time		creat(), write(), <b>close()</b>
2	First vs Subsequent access - Client fetches file that existed in the server		<b>1st open()</b> , close() 2nd open(), close()
3	Write vs Read cost		<b>read()</b> , <b>write()</b>
4	Scalability (open() and close() cost wrt multiple client processes on the <b>same</b> machine)	(i) Client creates the file for the first time	open(), <b>close()</b>
		(ii) Client fetches file that existed in the server	<b>open()</b> , close()
5	Experimental (open() and close() cost wrt gRPC packet size)	(i) Client creates the file for the first time	open(), <b>close()</b>
		(ii) Client fetches file that existed in the server	<b>open()</b> , close()

**\*Costly operations highlighted in bold**

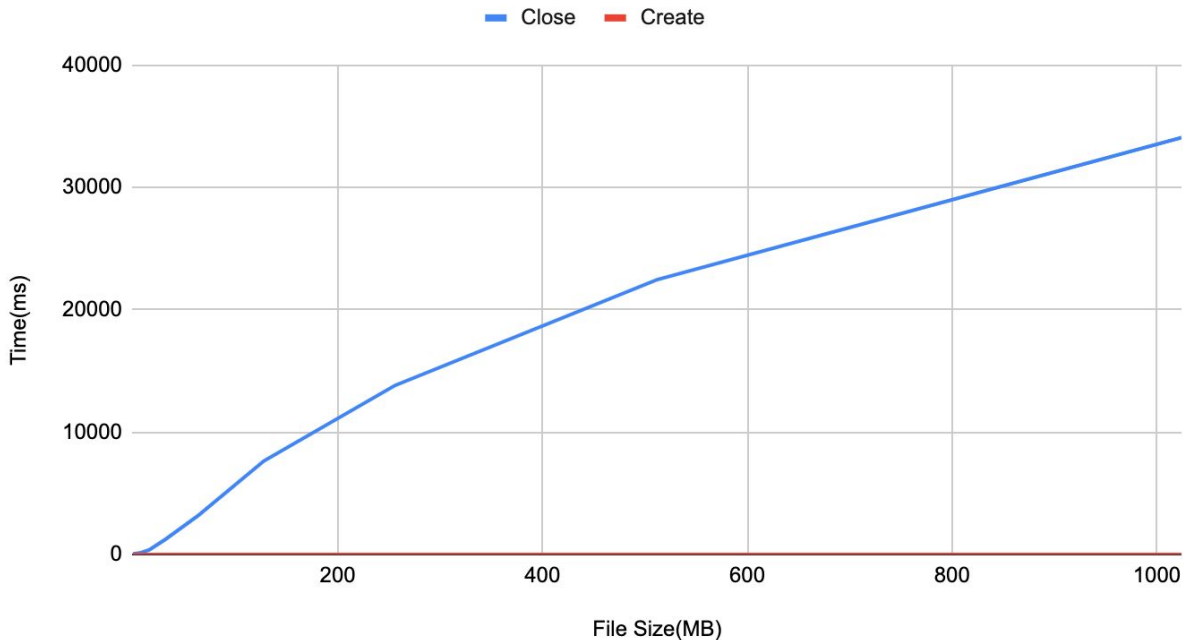
# Case - 1: Client creates the file for the first time



## Observations:

- Create time is independent to the file size
- Close time increases with increase in file size as expected

Client creating file for first time



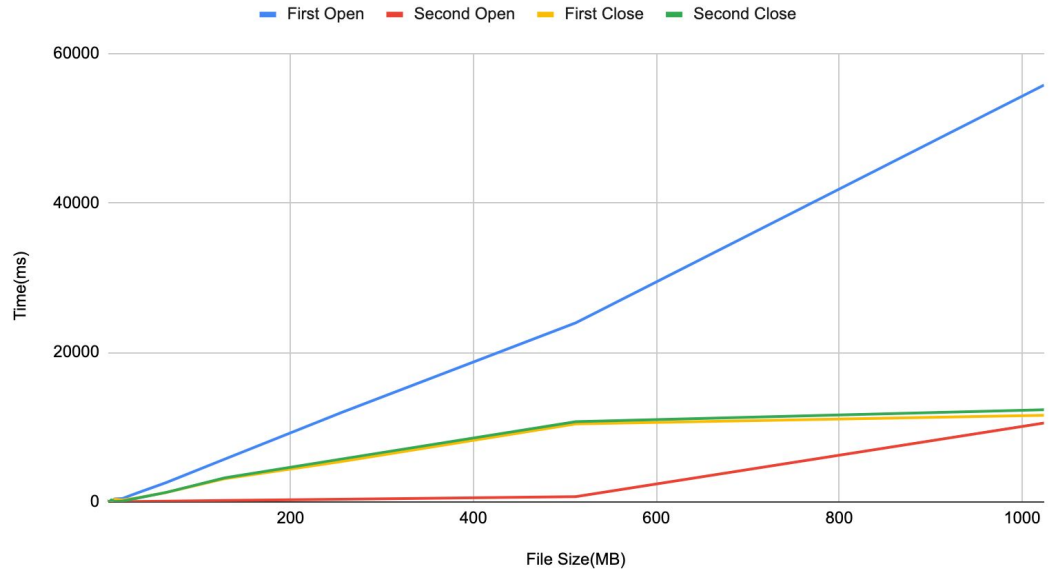


## Case - 2: Client fetches file that existed in the server

### Observations:

- Only first open is increasing with file size. Since, file is not changed by the client, other operations are not varying with file size

First Open, Second Open, First Close and Second Close




## Case - 3: Write vs Read

### Observations:

- No server/GRPC communication is involved here
- As expected, read and write time increases with file size.
- These are slower than normal read/write because of fuse intercept

Write vs Read



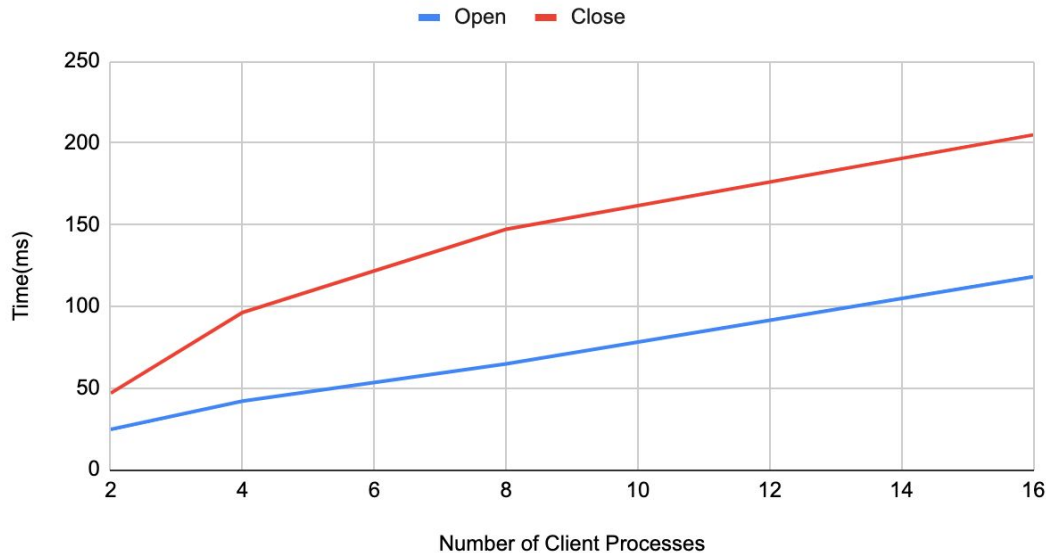
## Case - 4: Scalability

Multiple application processes are spawned to perform open and close operations involving gRPC

### Observations:

- Latency increases with increase in number of client processes
- This is due to the fact that client side fuse implementations are currently single threaded.

Scalability



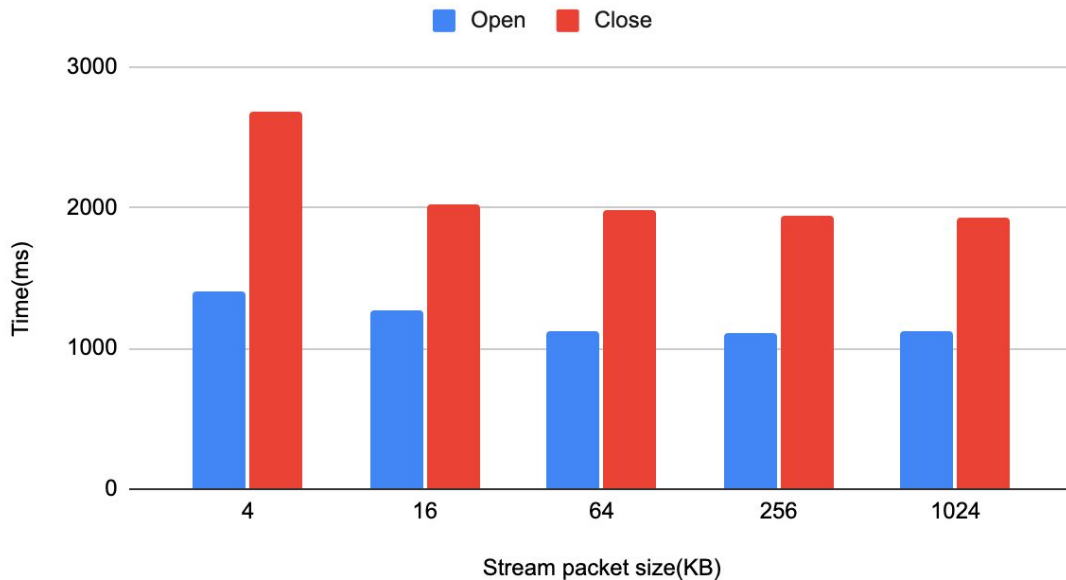
## Case - 5: Experimental

(cost wrt gRPC stream packet size)

### Observations:

- We experimented gRPC streams with different packet size during open() and close()
- Higher packet size improves open() and close() latency, and reaches optimal around 256KB

Performance on different stream packet size



# Performance Improvement Strategies



Only open() and close() calls involve network communication. So we focused on them for performance improvement.

1. Open() - getAttr to check last modified time, only fetch file if modified on server
2. Close() - Flush to server only if file was modified
3. Open/Close() - Streaming for concurrent reads and writes on client-server



Demos

# Demos



1. E2E File-System Demo
2. Last Writer Wins Demo
3. Client Crash Demo
4. Server Crash Demo
5. gRPC Retry Demo

  
**Thank you!**

